

Lab Sheet 2

Sabaragamuwa University of Sri Lanka

Faculty of Computing

Software Engineering

SE6103 - Parallel and Distributed Systems

Name : S.P.I.L.D. Samaraweera

Reg No : 19APSE4305

Academic Period : 3rd Year 2nd Semester

Due Date : 25/11/2024

Setting Up the Distributed Hadoop Cluster

```
version: '3'
services:
  namenode:
    image: bde2020/hadoop-namenode:latest
    container_name: namenode
    environment:
      - CLUSTER_NAME=testhadoop
      - CORE_CONF_fs_defaultFS=hdfs://namenode:8020
    ports:
      - "9870:9870" # Web UI
      - "9000:9000" # NameNode port
    volumes:
      - namenode-data:/hadoop/dfs/name

  datanode:
    image: bde2020/hadoop-datanode:latest
    container_name: datanode
    environment:
      - CORE_CONF_fs_defaultFS=hdfs://namenode:8020
    volumes:
      - datanode-data:/hadoop/dfs/data

  depends_on:
    - namenode

  historyserver:
    image: bde2020/hadoop-historyserver:latest
    container_name: historyserver
    depends_on:
      - namenode
      - datanode
    ports:
      - "8188:8188" # Job History server UI
    environment:
      - CORE_CONF_fs_defaultFS=hdfs://namenode:8020

volumes:
  namenode-data:
  datanode-data:
```

Deploy the Cluster

```
PS C:\Docker> docker-compose up -d
[+] Running 18/5
 ✓ historyserver 9 layers [██████████] 0B/0B Pulled
 ✓ datanode 7 layers [██████████] 0B/0B Pulled

[+] Running 5/5
[+] Running 5/6er_default Created
[+] Running 5/6er_default Created
[+] Running 5/6er_default Created
[+] Running 5/6er_default Created
[+] Running 5/6er_default Created
[+] Running 5/6er_default Created
[+] Running 5/6er_default Created
```

```
● PS C:\Docker> docker exec -it namenode hdfs dfs -mkdir -p /input
```

Verify Cluster Status

[Hadoop](#) [Overview](#) [Datanodes](#) [Datanode Volume Failures](#) [Snapshot](#) [Startup Progress](#) [Utilities](#)

Overview 'namenode:8020' (active)

Started:	Mon Nov 25 14:10:14 +0530 2024
Version:	3.2.1, rb3cbbb467e22ea829b3808f4b7b01d07e0bf3842
Compiled:	Tue Sep 10 21:26:00 +0530 2019 by rohithsharmaks from branch-3.2.1
Cluster ID:	CID-e2add280-0c9c-432b-82c3-0454ba541b99
Block Pool ID:	BP-423021675-172.19.0.2-1732524011563

Download Sample Data

Upload Data to the HDFS

```
● PS C:\Docker> docker exec -it namenode hdfs dfs -put /sample.txt /input
2024-11-25 08:19:46,116 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
```

```
● PS C:\Docker> docker exec -it namenode ls /sample.txt
/sample.txt
PS C:\Docker> docker exec -it namenode hadoop jar
```

Running a MapReduce Job

```
PS C:\Docker> docker exec -it namenode bash
root@61e13ba3012a:/# hadoop jar /opt/hadoop/share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar wordcount /input /output
```

View Job Output

```

root@61e13ba3012a:/# hadoop jar /opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.1.jar wordcount /input /output
2024-11-25 08:31:05,217 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2024-11-25 08:31:05,424 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2024-11-25 08:31:05,424 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2024-11-25 08:31:05,790 INFO input.FileInputFormat: Total input files to process : 1
2024-11-25 08:31:05,817 INFO mapreduce.JobSubmitter: number of splits:1
2024-11-25 08:31:06,120 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1280096721_0001

```

```

root@61e13ba3012a:/# exit
exit
● PS C:\Docker> docker exec -it namenode hdfs dfs -ls /output
Found 2 items
-rw-r--r--  3 root supergroup          0 2024-11-25 08:31 /output/_SUCCESS
-rw-r--r--  3 root supergroup    124 2024-11-25 08:31 /output/part-r-00000

```

Analyze and Cleanup

```

● PS C:\Docker> docker exec -it namenode hdfs dfs -cat /output/part-r-00000
2024-11-25 08:33:48,085 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
"Hadoop 1
> 1
a 1
and 1
big 1
data." 1
distributed 1
echo 1
for 1
framework 1
is 1
of 1
processing 1
framework 1
is 1
○ framework 1
is 1
of 1
framework 1
is 1
framework 1
framework 1
is 1
of 1
processing 1
sample.txt 1
storage 1

```

```

● PS C:\Docker> docker-compose down
[+] Running 4/4
✓ Container historyserver Removed
✓ Container datanode      Removed
✓ Container namenode      Removed
✓ Network docker_default  Removed

```